

Paul Williams

I specialize in knowledge engineering, content architecture, tooling, and automation for enterprise software content teams, with an emphasis on AI-ready knowledge pipelines and scalable content systems. Domain expertise includes cloud-native computing, developer documentation, and scalable content systems.

PROFESSIONAL EXPERIENCE

Senior DITA Information Architect, Splunk, a Cisco company

JANUARY 2025 - PRESENT

- Built an AI authoring platform that combines retrieval over internal documentation, deterministic DITA XML validation, and reusable authoring workflows to transform Markdown drafts into production-ready DITA content, accelerating structured authoring while maintaining quality.
- Architected a nightly documentation pipeline that exports approximately 100,000 CCMS topics into an AI-optimized Markdown repository with a rich metadata index, enabling efficient retrieval by LLM-powered tools and agentic workflows.
- Designed and implemented an automated PDF publishing platform that keeps documentation synchronized with product releases through a nightly CI pipeline, incorporating observability, retry logic, state tracking, and resilient multi-worker processing.
- Led the migration of two enterprise documentation sets from legacy platforms to a DITA-based CCMS, designing Python conversion pipelines that transformed tens of thousands of topics while preserving metadata, cross references, and document structure with minimal manual revision.
- Established observability and operational tooling, including metadata quality reporting, deployment notifications, GitLab synchronization, and advanced XPath search capabilities that improve visibility into documentation operations.
- Advanced content governance by defining sitemap policies, metadata standards, resource ID architecture, conditional processing, and reuse practices that improve release reliability, publishing consistency, and long-term maintainability.
- Partnered with writers, editors, and software engineers to define documentation architecture, drive adoption of AI-assisted authoring, deliver technical enablement, and evolve the organization's documentation platform.

Staff Technical Writer, Splunk, a Cisco company

OCTOBER 2023 - JANUARY 2025

- Created a new release notes section in the Splunk Observability Cloud documentation, collaborating with stakeholders in Product, Engineering, Marketing, and Product Experience.
- Wrote and maintained Splunk Observability Cloud instrumentation documentation across eight programming languages, leveraging CI/CD pipelines to achieve consistency between engineering repositories and comprehensive reference documentation.
- Co-led a documentation team focus group on using generative AI in technical communications, spearheading discussions on optimizing documentation for LLM consumption and leveraging AI to enhance Information Experience team productivity.
- Continually improved information architecture, content discoverability, and user experience of the Splunk Observability Cloud developer portal, including the Developer Guide and API Reference.
- Contributed to open-source OpenTelemetry documentation, aligning downstream terminological choices with upstream usage trends.
- Wrote internal documentation for the Splunk Cloud operations teams, contributing to a significant decrease in the remediation of disk space usage incidents and leveraging Python and GitLab to create a monthly runbook metrics report.

Technical Writer, Mirantis

Portland, OR | pauljwil@gmail.com | Portfolio: pauljwil.com | linkedin.com/in/pauljwil | https://github.com/pauljwil

OCTOBER 2020 - OCTOBER 2023

- Contributed to and maintained approximately three thousand pages of technical documentation across multiple containerization, orchestration, and image registry products, using rST, Sphinx, Gerrit, and Jira, reducing Customer Support tickets and empowering internal teams to effectively deploy and use Mirantis Containers software.
- Participated in documentation team content strategy decisions, advocating for the creation of quick start guides, best practices content, user tutorials, and site architecture reorganization, and led efforts to set documentation standards, processes, and best practices.
- Conducted thorough testing of documentation changes and new product features using AWS EC2 instances, Kubernetes, Helm, Docker Desktop, and internal Mirantis tooling.
- Created and maintained YAML code samples used in the management of Kubernetes clusters.
- Automated the editing process of lengthy and repetitive documentation changes by developing Python and Bash scripts, improving efficiency and accuracy.
- Revitalized the API documentation for two products by migrating from Swagger to Redoc to enhance the overall presentation.
- Analyzed customer feedback to enhance product documentation at multiple stages in the release process.
- Successfully restructured a comprehensive documentation repository inherited from a company merger. Ensured adherence to industry-standard IBM Style Guide and Mirantis documentation team best practices, improving consistency, clarity, and usability of the documentation.

Software Developer, Mirantis

MARCH 2022 - OCTOBER 2023 (PART-TIME)

- Developed a Prometheus exporter for Mirantis Secure Registry, leveraging extensive querying of RethinkDB to provide robust monitoring capabilities.
- Implemented a telemetry feature for Mirantis Migration Tool, enabling the collection and analysis of usage metrics using segment.com.
- Resolved bugs in two Mirantis products, including the introduction of context into RethinkDB functions, logging enhancements, and fixing HTTP errors in MSR endpoints.
- Conducted thorough testing of software changes using AWS EC2 instances, Kubernetes, Helm, Docker Desktop, and internal Mirantis tooling.
- Utilized strong writing skills to contribute to internal documentation, ERDs, design docs, tickets, and code comments, ensuring clear and comprehensive project documentation.
- Actively participated in Agile processes, including daily standups, sprint planning, retrospective, ticket grooming, and collaborative troubleshooting sessions.

SKILLS AND TOOLS

- **Knowledge Engineering:** AI Knowledge Pipelines, RAG Systems, MCP Servers, Agentic Workflows
- **Content Engineering:** Python, Bash, XSLT, DITA Open Toolkit, JavaScript, CSS
- **Structured Authoring:** DITA XML, Heretto, Metadata Governance, Content Strategy
- **Docs as Code:** Sphinx, reStructuredText, Git, GitHub, GitLab, OpenAPI
- **DevOps/Cloud:** Go, Docker, Kubernetes, Helm, AWS

EDUCATION

- **Doctoral studies, Philosophy,** University of Texas at Austin
- **Master of Arts, Philosophy,** San Francisco State University
- **Bachelor of Music, Composition,** San Francisco State University